

# CS 5/7320 Artificial Intelligence

## Intelligent Agents AIMA Chapter 2

---

Slides by Michael Hahsler  
based on slides by Svetlana Lazepnik  
with figures from the AIMA textbook.



This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-sa/4.0/).

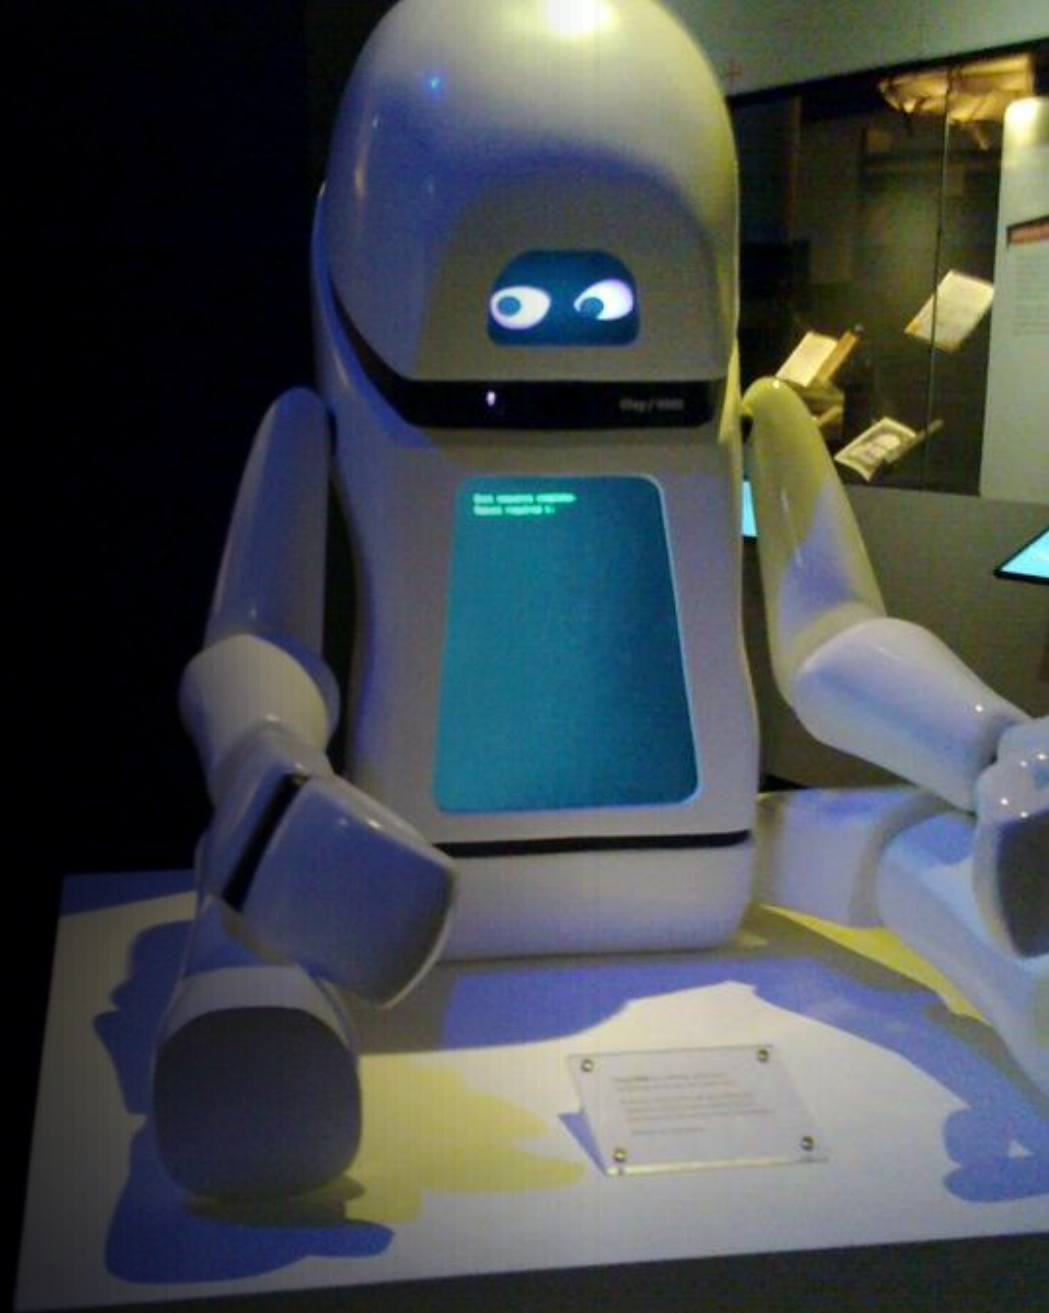
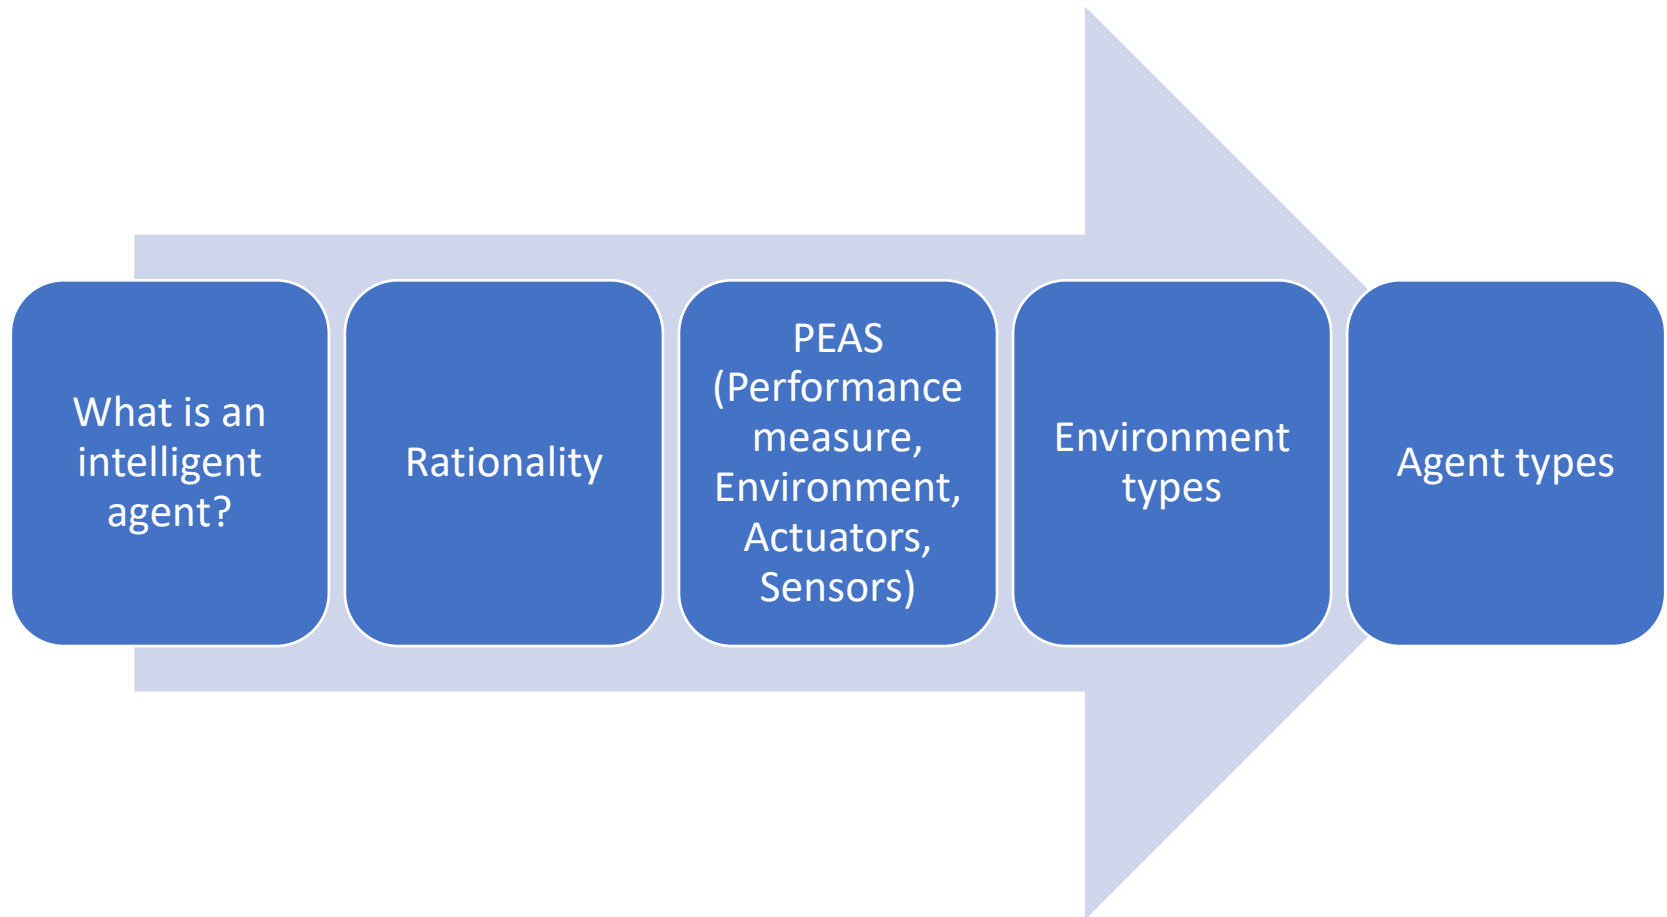
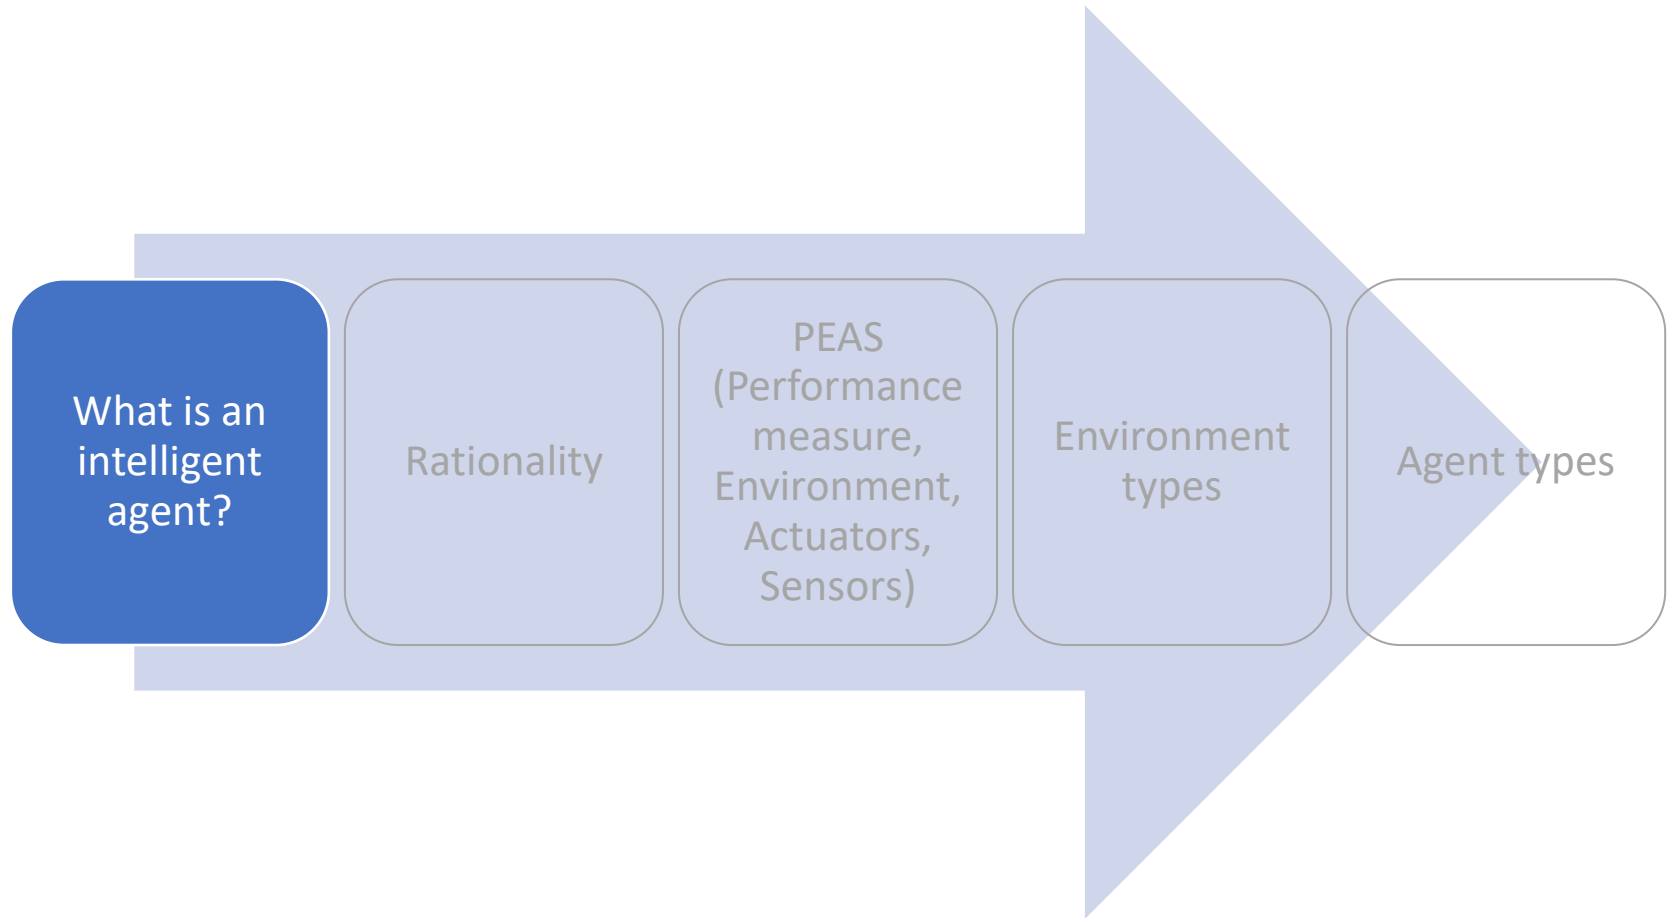


Image: "Robot at the British Library Science Fiction Exhibition"  
by BadgerGravling

# Outline

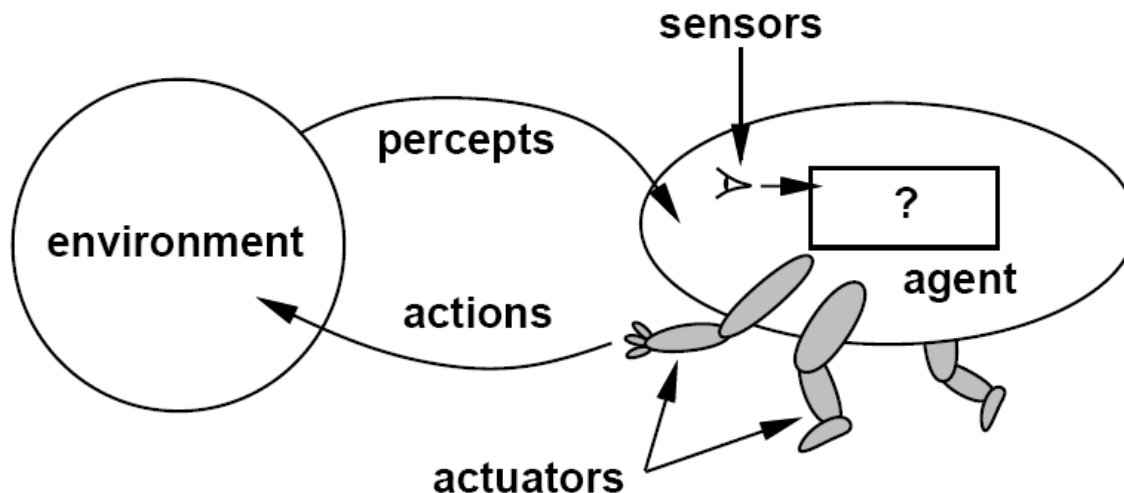


# Outline



# What is an Agents?

- An **agent** is anything that can be viewed as **perceiving** its **environment** through **sensors** and **acting** upon that environment through **actuators**.

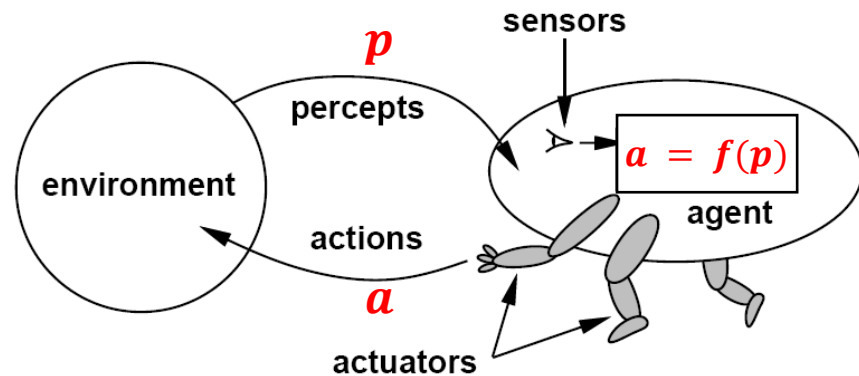


- **Control theory:** A **closed-loop control system** (= feedback control system) is a set of mechanical or electronic devices that automatically regulate a process variable to a desired state or set point without human interaction. The agent is called a controller.
- **Softbot:** Agent is a software program that runs on a host device.

# Agent Function and Agent Program

The **agent function** maps from the set of all possible *percept sequences*  $P^*$  to the *set of actions*  $A$  formulated as an abstract mathematical function.

$$f : P^* \rightarrow A$$



The **agent program** is a concrete implementation of this function for a given physical system.

Agent = architecture (hardware) + agent program (implementation of  $f$ )



- Sensors
- Memory
- Computational power

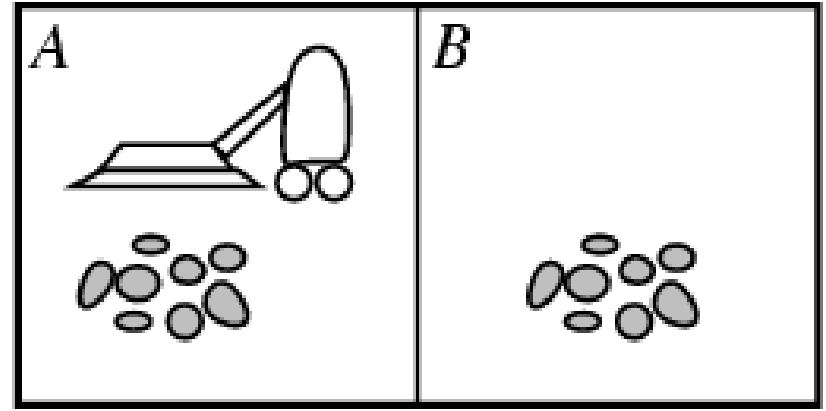
# Example: Vacuum-cleaner World

- **Percepts:**

Location and status,  
e.g., [A, Dirty]

- **Actions:**

Left, Right, Suck, NoOp



Most recent  
Percept  $p$

Agent function:  $f : P^* \rightarrow A$

Percept Sequence	Action
[A, Clean]	Right
[A, Dirty]	Suck
...	
[A, Clean], [B, Clean]	Left
...	
[A, Clean], [B, Clean], [A, Dirty]	Suck
...	

Implemented agent program:

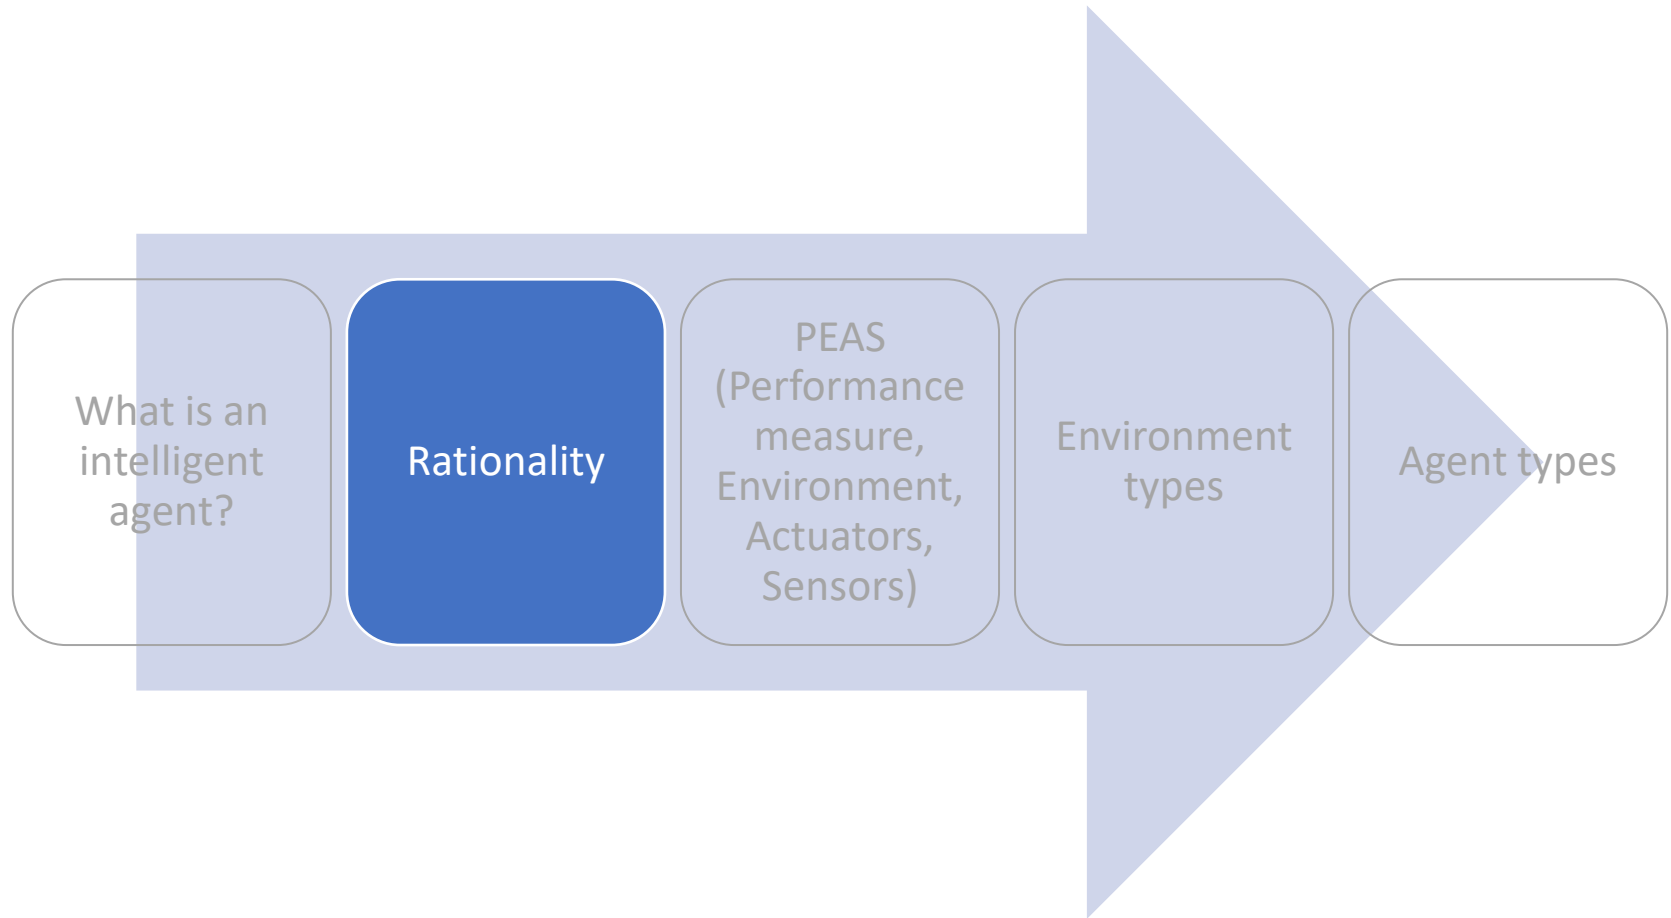
**function Vacuum-Agent**([location, status])

returns an **action  $a$**

```
if status = Dirty then return Suck
else if location = A then return Right
else if location = B then return Left
```

**Problem:** This table can become infinitively large!

# Outline



# Rational Agents: What is Good Behavior?

## Foundation

- **Consequentialism:** Evaluate behavior by its consequences.
- **Utilitarianism:** Maximize happiness and well-being.

## Definition of a rational agent:

*“For each possible percept sequence, a rational agent should select an **action** that **maximizes its expected performance measure**, given the evidence provided by the **percept sequence** and the **agent’s built-in knowledge**.”*

- **Performance measure:** An *objective* criterion for success of an agent's behavior (often called utility function or reward function).
- **Expectation:** Outcome averaged over all possible situations that may arise.

## This means:

- **Rationality is only an ideal**
- **Rationality  $\neq$  Omniscience** (rational agents can make mistakes if percepts and knowledge do not suffice to make a good decision)
- **Rationality  $\neq$  Perfection** (rational agents maximize **expected** outcomes not actual outcomes)
- **It is rational to explore and learn** (i.e., use **percepts** to supplement prior knowledge and become autonomous)



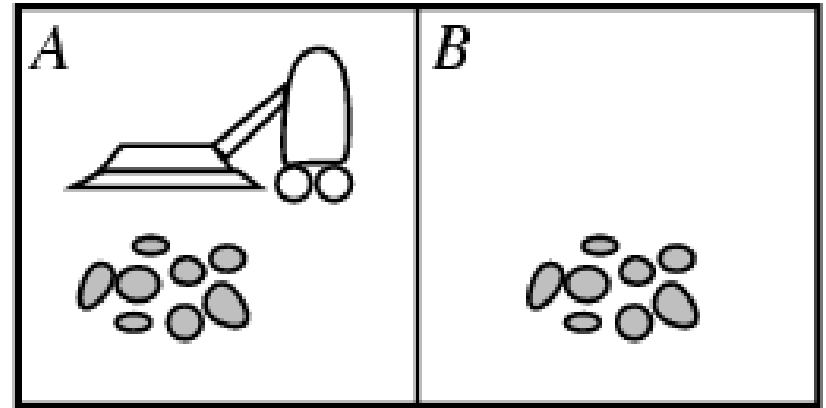
# Example: Vacuum-cleaner World

- **Percepts:**

Location and status,  
e.g., [A, Dirty]

- **Actions:**

Left, Right, Suck, NoOp



Agent function:

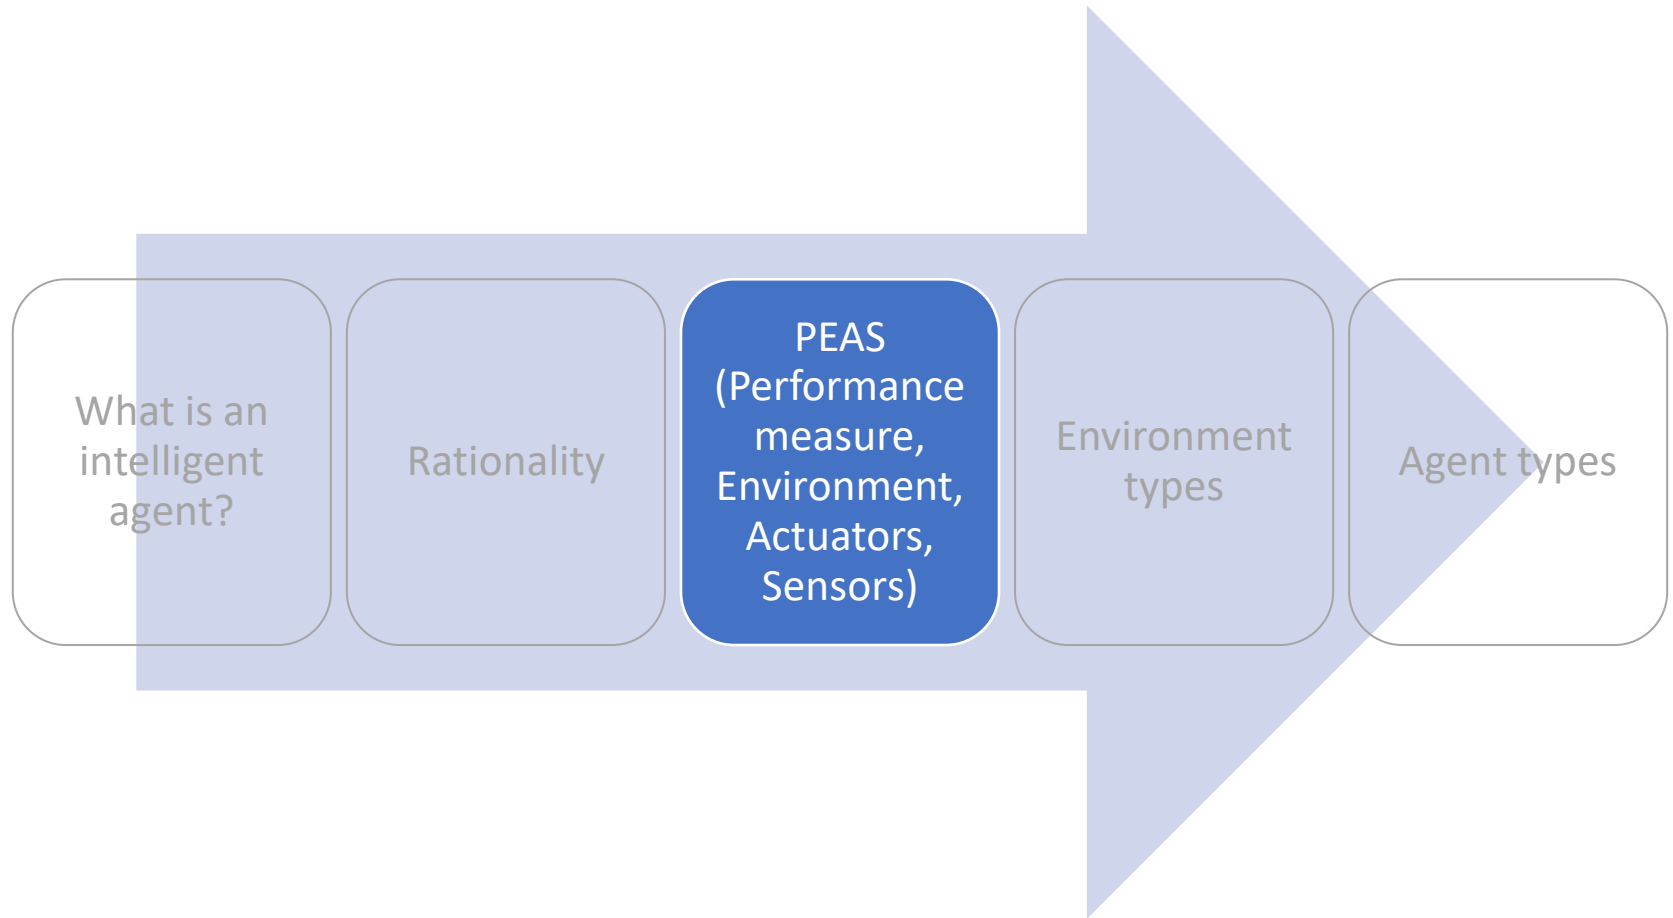
Percept Sequence	Action
[A, Clean]	Right
[A, Dirty]	Suck
...	
[A, Clean], [B, Clean]	Left
...	

Implemented agent program:

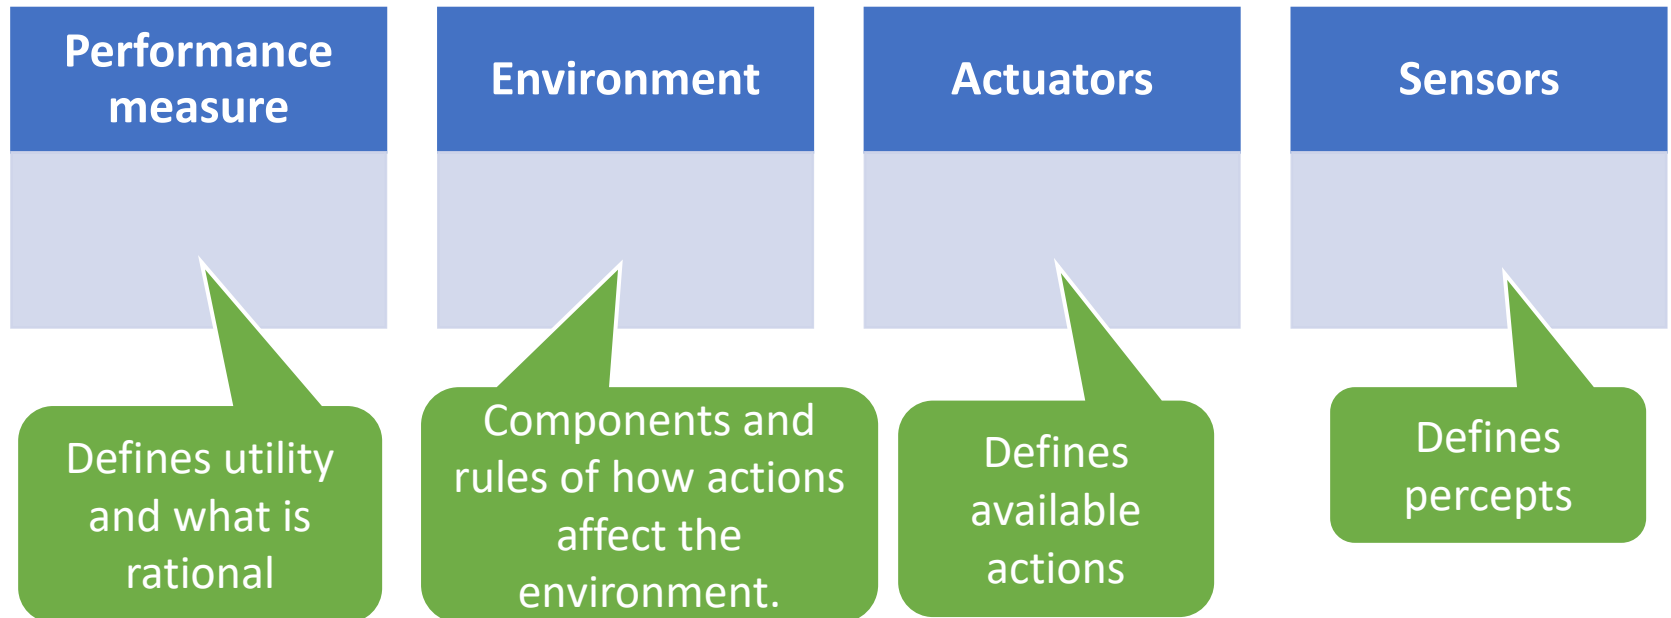
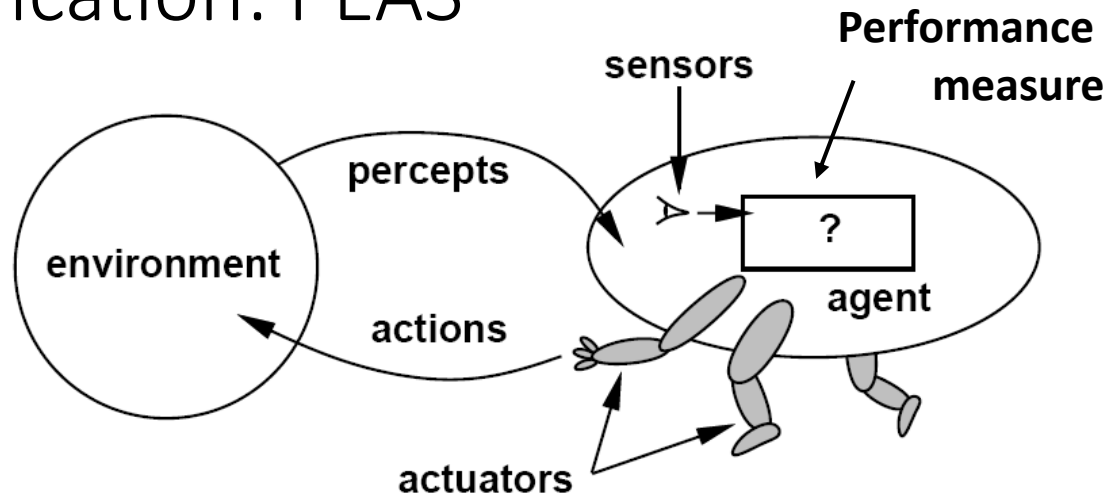
```
function Vacuum-Agent([location, status])  
    returns an action  
  
    if status = Dirty then return Suck  
    else if location = A then return Right  
    else if location = B then return Left
```

**What could be a performance measure?**  
**Is this agent program rational?**

# Outline



# Problem Specification: PEAS



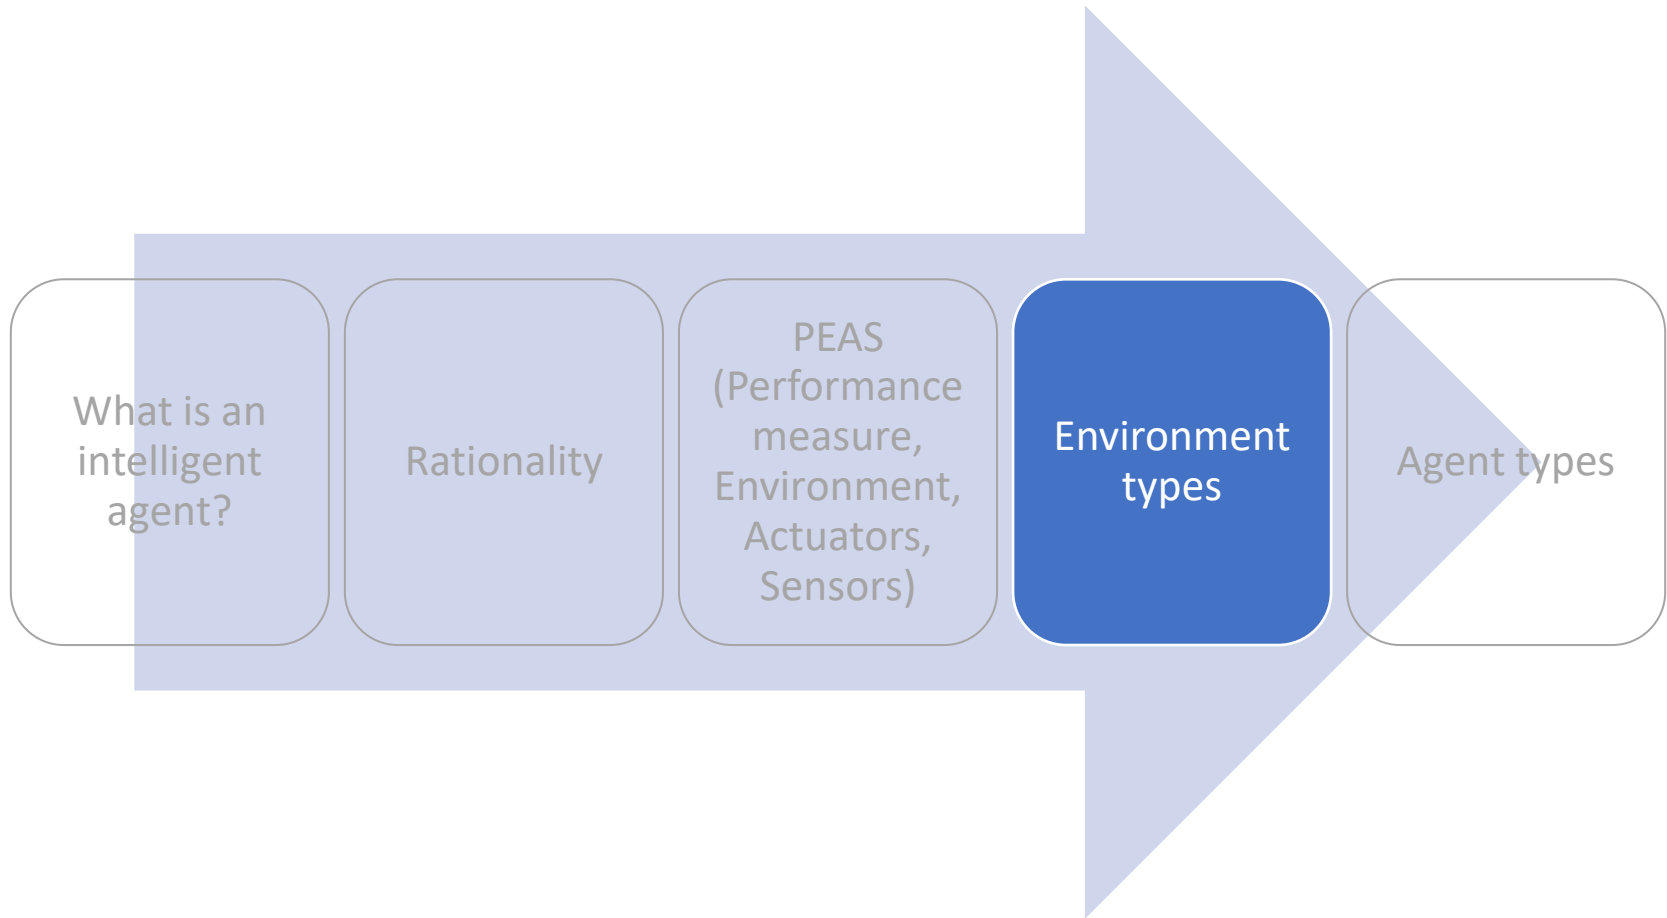
# Example: Automated Taxi Driver

Performance measure	Environment	Actuators	Sensors
<ul style="list-style-type: none"><li>• Safe</li><li>• fast</li><li>• legal</li><li>• comfortable trip</li><li>• maximize profits</li></ul>	<ul style="list-style-type: none"><li>• Roads</li><li>• other traffic</li><li>• pedestrians</li><li>• customers</li></ul>	<ul style="list-style-type: none"><li>• Steering wheel</li><li>• accelerator</li><li>• brake</li><li>• signal</li><li>• horn</li></ul>	<ul style="list-style-type: none"><li>• Cameras</li><li>• sonar</li><li>• speedometer</li><li>• GPS</li><li>• Odometer</li><li>• engine sensors</li><li>• keyboard</li></ul>

# Example: Spam Filter

Performance measure	Environment	Actuators	Sensors
<ul style="list-style-type: none"><li>• Accuracy: Minimizing false positives, false negatives</li></ul>	<ul style="list-style-type: none"><li>• A user's email account</li><li>• email server</li></ul>	<ul style="list-style-type: none"><li>• Mark as spam</li><li>• delete</li><li>• etc.</li></ul>	<ul style="list-style-type: none"><li>• Incoming messages</li><li>• other information about user's account</li></ul>

# Outline



# Environment Types

**Fully observable:** The agent's sensors give it access to the complete state of the environment. The agent can “see” the whole environment.

**Deterministic:** The next state of the environment is completely determined by the current state and the agent's action.

**Known:** The agent knows the rules of the environment and can predict the outcome of actions. It knows the transition function.

**vs. Partially observable):** The agent cannot see all aspects of the state. E.g., it can't see through walls.

**vs. Stochastic:** The next state cannot be determined from the current state and the action (there is some randomness).

**Strategic:** The environment is adversarial and chooses actions strategically to harm the agent. E.g., a game where the other player is modeled as part of the environment.

**vs. Unknown:** The agent cannot predict the outcome of actions. It needs to learn the transition function by trying actions.

# Environment Types

**Static:** The environment is **not** changing while agent is deliberating.

**Semidynamic:** the environment does not change while deliberating, but the agent's performance score depends on how fast it acts.

**Discrete:** The environment provides a fixed number of distinct percepts, actions, and environment states. Time can also evolve in a discrete or continuous fashion.

**Episodic:** Episode = a self-contained sequence of actions. The agent's choice of action in one episode does not affect the next episodes. The agent does the same task repeatedly.

**Single agent:** An agent operating by itself in an environment.

**vs. Dynamic:** The environment is changing while the agent is deliberating.

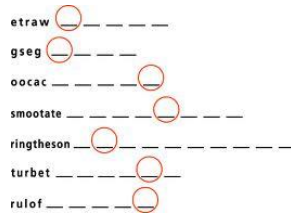
**vs. Continuous:** Percepts, actions, state variables or time are continuous leading to an infinite state, percept or action space.

**vs. Sequential:** Actions now affect the outcomes later. E.g., learning makes problems sequential.

**vs. Multi-agent:** Agent cooperate or compete in the same environment.



# Examples of Different Environments



Word jumble solver



Chess with a clock



Scrabble

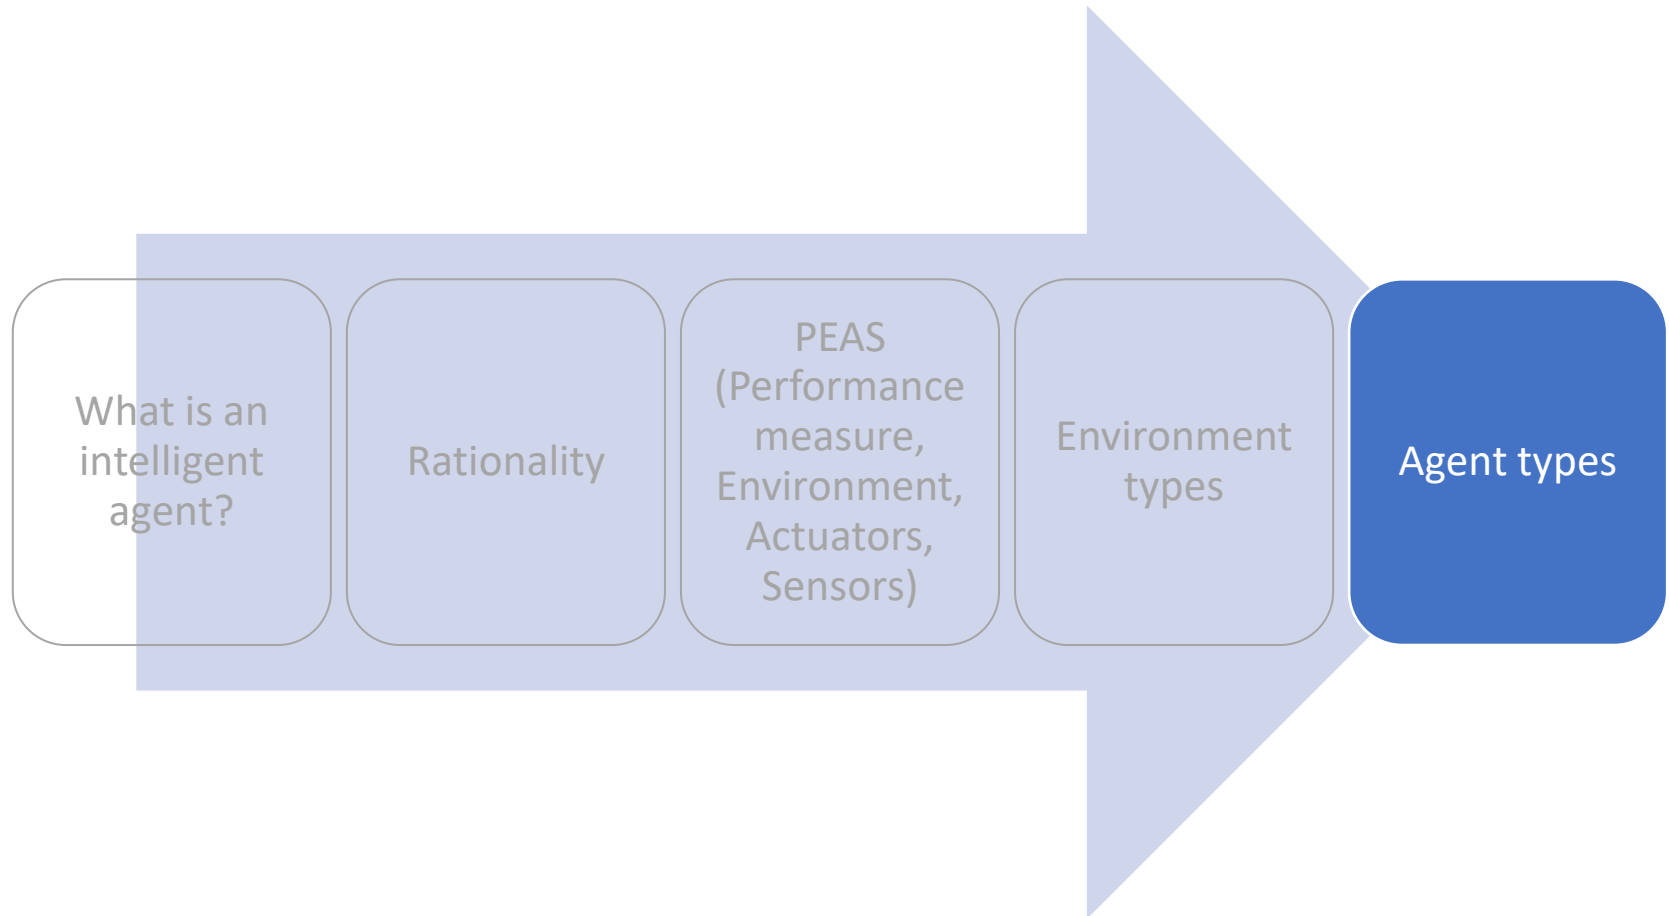


Taxi driving

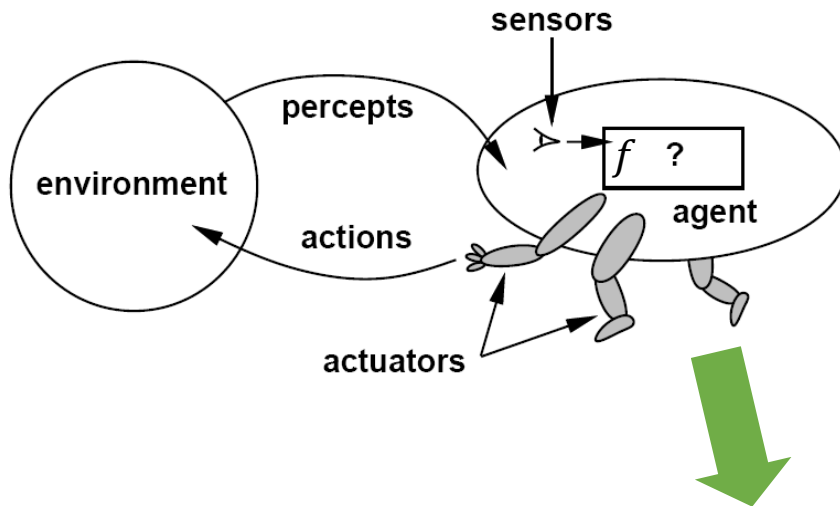
Observable	Fully	Fully	Partially	Partially
Deterministic	Deterministic	Strategic	Stochastic +Strategic	Stochastic
Episodic?	Episodic	Episodic	Episodic	Sequential
Static	Static	Semidynamic	Static	Dynamic
Discrete	Discrete	Discrete	Discrete	Continuous
Single agent	Single	Multi*	Multi*	Multi*

\* Can be models as a single agent problem with the other agent(s) in the environment.

# Outline



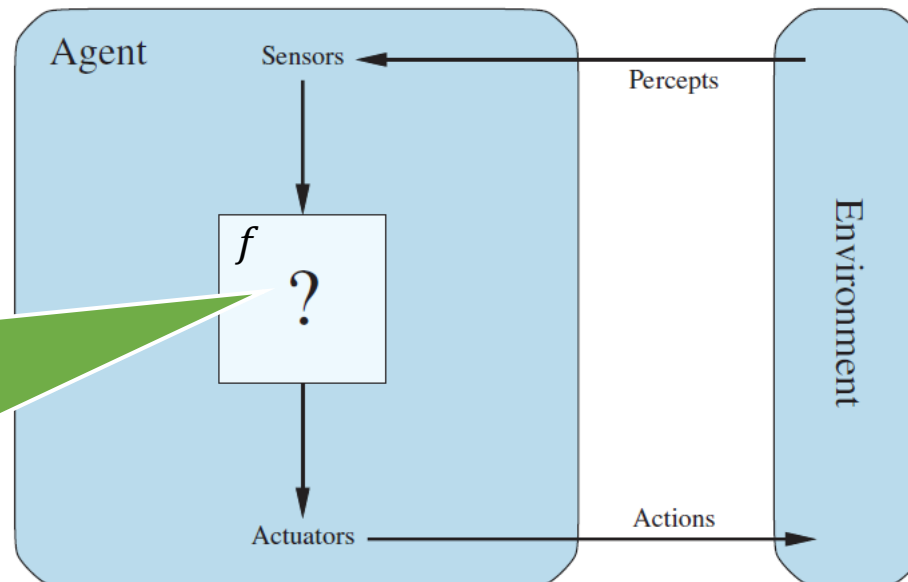
# Designing a Rational Agent



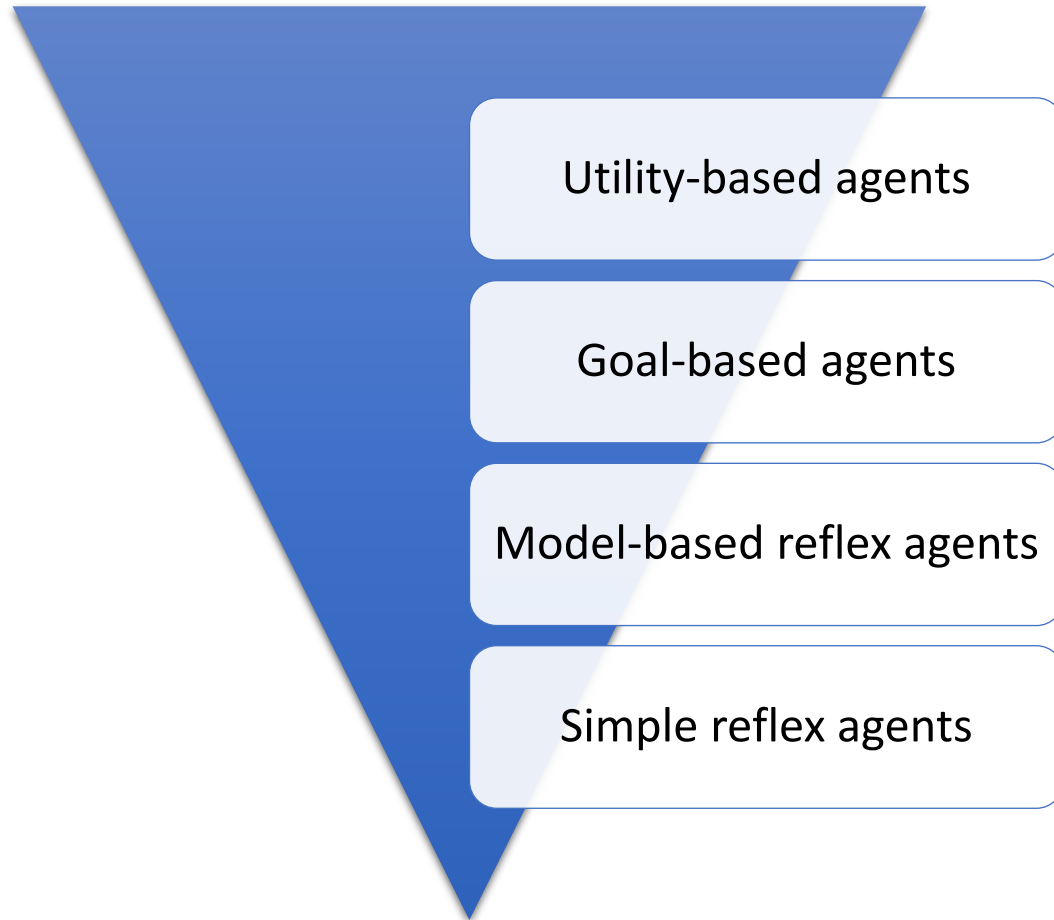
Remember the definition of a rational agent:

*“For each possible percept sequence, a rational agent should select an **action** that **maximizes its expected performance measure**, given the evidence provided by the **percept sequence** and the **agent’s built-in knowledge**.”*

- Assess performance measure
- Remember percept sequence
- Built-in knowledge

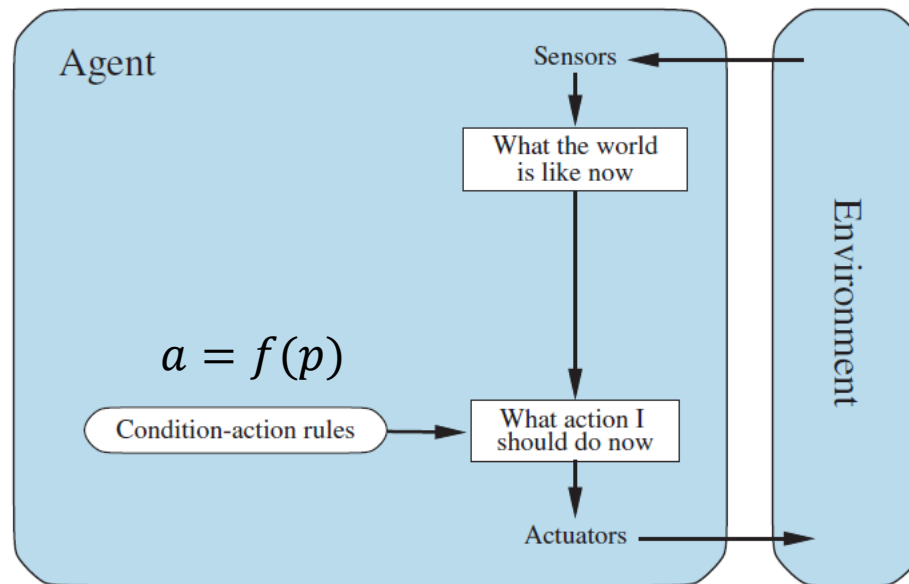


# Hierarchy of Agent Types



# Simple Reflex Agent

- Uses only built-in knowledge in the form of **rules** that select action only **based on the current percept**. This is typically very fast!
- The **agent does not know about the performance measure**! But well-designed rules can lead to good performance.
- The agent needs no memory and ignores all past percepts.

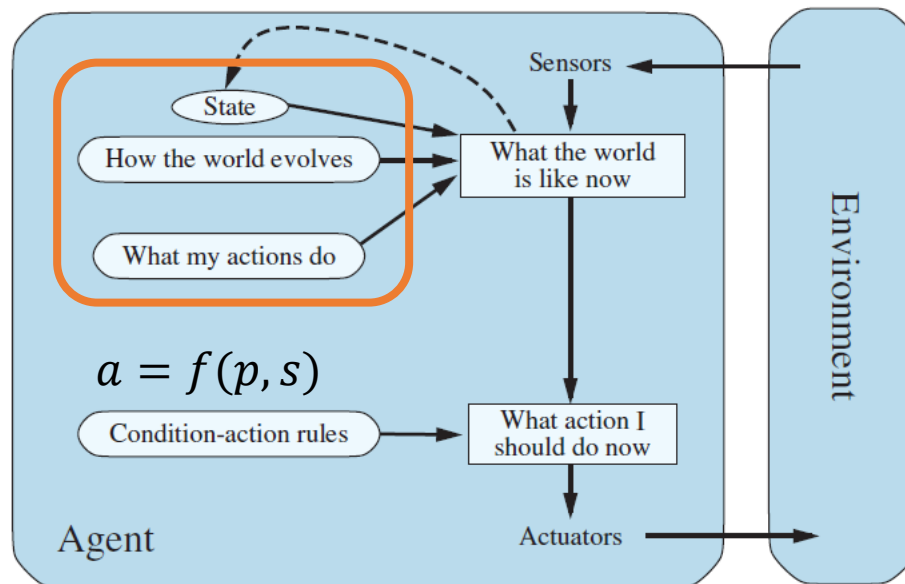


The interactions is a sequence:  $p_0, a_0, p_1, a_1, p_2, a_2, \dots, p_t, a_t, \dots$

**Example:** A simple vacuum cleaner that uses rules based on its current sensor input.

# Model-based Reflex Agent

- Maintains a **state variable** to keep track of aspects of the environment that cannot be currently observed. I.e., it has memory and knows how the environment reacts to actions.
- The state is updated using the percept.
- There is now more information for the **rules** to make better decisions.



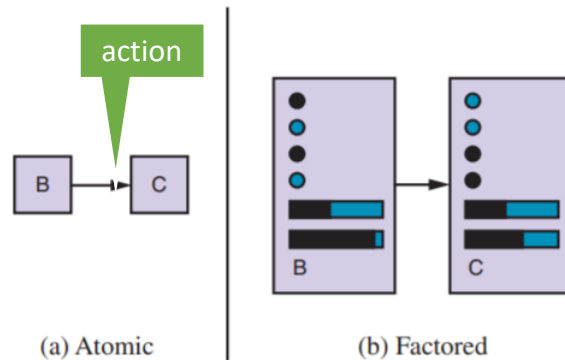
The interactions is a sequence:  $s_0, a_0, p_0, s_1, a_1, p_1, s_2, a_2, p_2, \dots, s_t, a_t, p_t, \dots$

**Example:** A vacuum cleaner that remembers where it has already cleaned.

# State Representation

States help to keep track of the environment. The representation can be

- **Atomic:** Just a label for a black box. E.g., A, B
- **Factored:** A vector of attribute values. E.g., [location = left, status = clean, temperature = 75 deg. F]



Actions cause a transition from one state to another.

We often construct atomic labels from factored information. E.g.: If the agent's factored state is the coordinate  $x = 7$  and  $y = 3$ , then the atomic state label could be the string  $(7, 3)$ . With the label, we can only compare if two labels are the same. With the factored state we can calculate distances!

**State Space:** The set of all possible states  $S$ . This set is typically very large!

# Old-school vs. Smart Thermostat



Old-school thermostat

**Percepts**

**States**



Smart thermostat

**Percepts**

**States**



# Old-school vs. Smart Thermostat



Set  
temperature  
range



Change  
temperature  
when you  
are too  
cold/warm.

## Old-school thermostat

### Percepts

temperature:  
Low, ok, high

### States

No states need

## Smart thermostat

### Percepts

- Temp: deg. F
- Outside temp.
- Weather report
- Energy curtailment
- Someone walking by
- Someone changes temp.
- Day & time
- ...

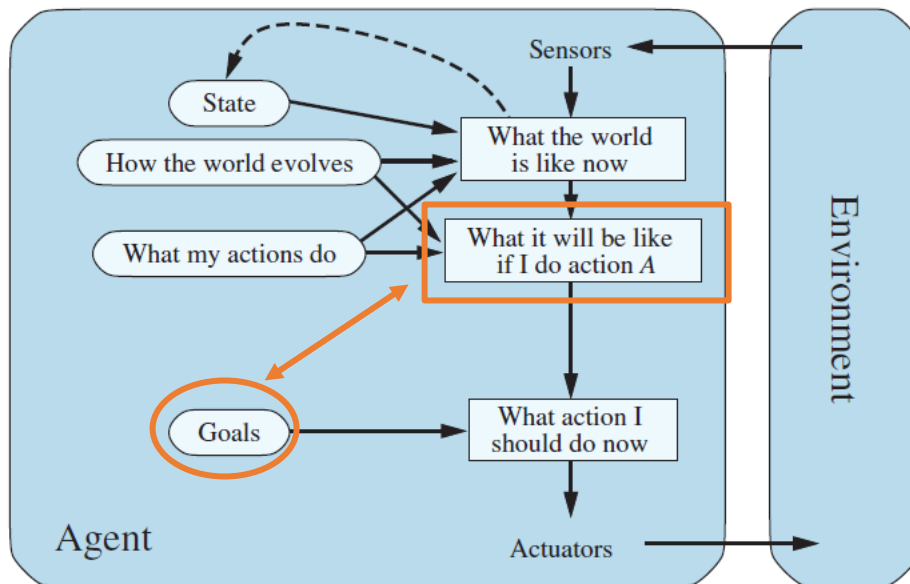
### States

Factored states

- Estimated time to cool the house
- Someone home?
- How long till someone is coming home?
- A/C: on, off

# Goal-based Agent

- The agent has the task to reach a defined **goal state** and is then finished.
- The agent needs to move towards the goal. It can use **search algorithms** to plan actions that lead to the goal.
- The performance measure is typically the cost to reach the goal.



$$a = \operatorname{argmin}_{a_0 \in A} \left[ \sum_{t=0}^T c_t \mid s_T \in S^{goal} \right]$$

Sum of the cost of a planned sequence of actions that leads to a goal state

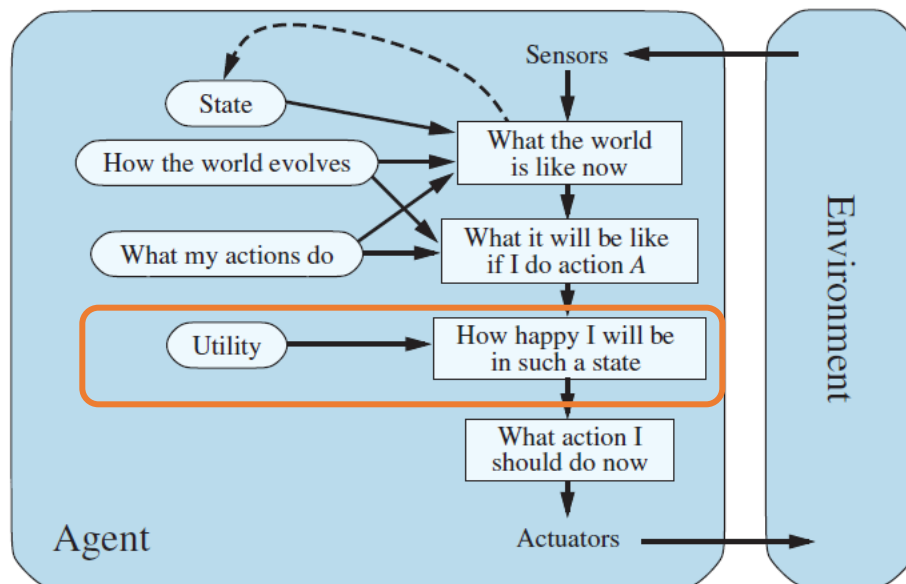
The interactions is a sequence:  $s_0, a_0, p_0, s_1, a_1, p_1, s_2, a_2, p_2, \dots, s^{goal}$

cost

**Example:** Solving a puzzle. What action gets me closer to the solution?

# Utility-based Agent

- The agent uses a utility function to evaluate the **desirability of each possible states**. This is typically expressed as the reward of being in a state  $R(s)$ .
- Choose actions to stay in desirable states.
- Performance measure: The discounted sum of expected utility over time.



$$a = \operatorname{argmax}_{a_0 \in A} \mathbb{E} \left[ \underbrace{\sum_{t=0}^{\infty} \gamma^t r_t}_{\text{Expected future discounted reward}} \right]$$

**Techniques:** Markov decision processes, reinforcement learning

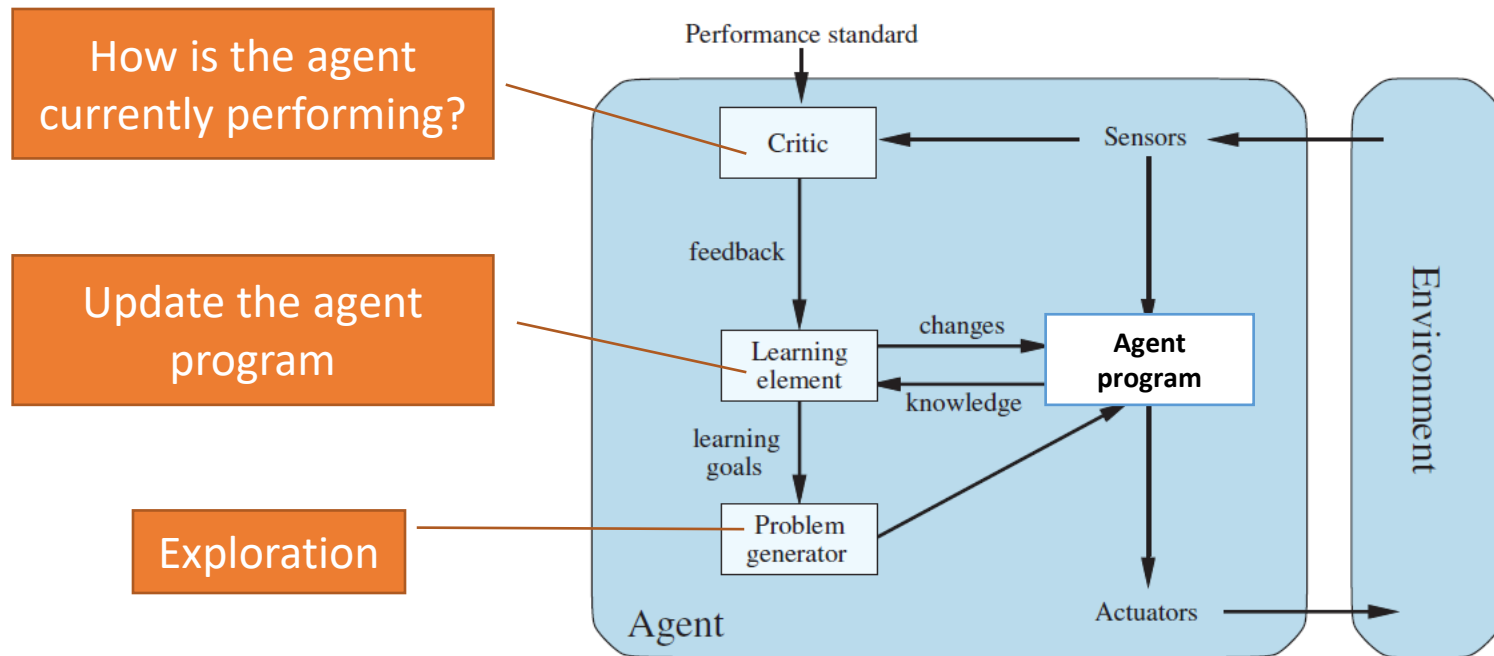
The interactions is a sequence:  $s_0, a_0, p_0, s_1, a_1, p_1, s_2, a_2, p_2, \dots$

$\underbrace{\hspace{10em}}_{\text{reward}}$

**Example:** An autonomous Mars rover prefers states where its battery is not critically low.

# Agents that Learn

The **learning element** modifies the agent program (reflex-based, goal-based, or utility-based) to improve its performance.



# Smart Thermostat



Change temperature when you are too cold/warm.

Reflex Agent?

Goal-based?

Utility-based?

## Smart thermostat

### Percepts

- Temp: deg. F
- Outside temp.
- Weather report
- Energy curtailment
- Someone walking by
- Someone changes temp.
- Day & time
- ...

### States

#### Factored states

- Estimated time to cool the house
- Someone home?
- How long till someone is coming home?
- A/C: on, off

# What Type of Intelligent Agent is this?

Features are:

- Control via App
- Cleaning Modes
- Navigation
- Mapping
- Boundary blockers



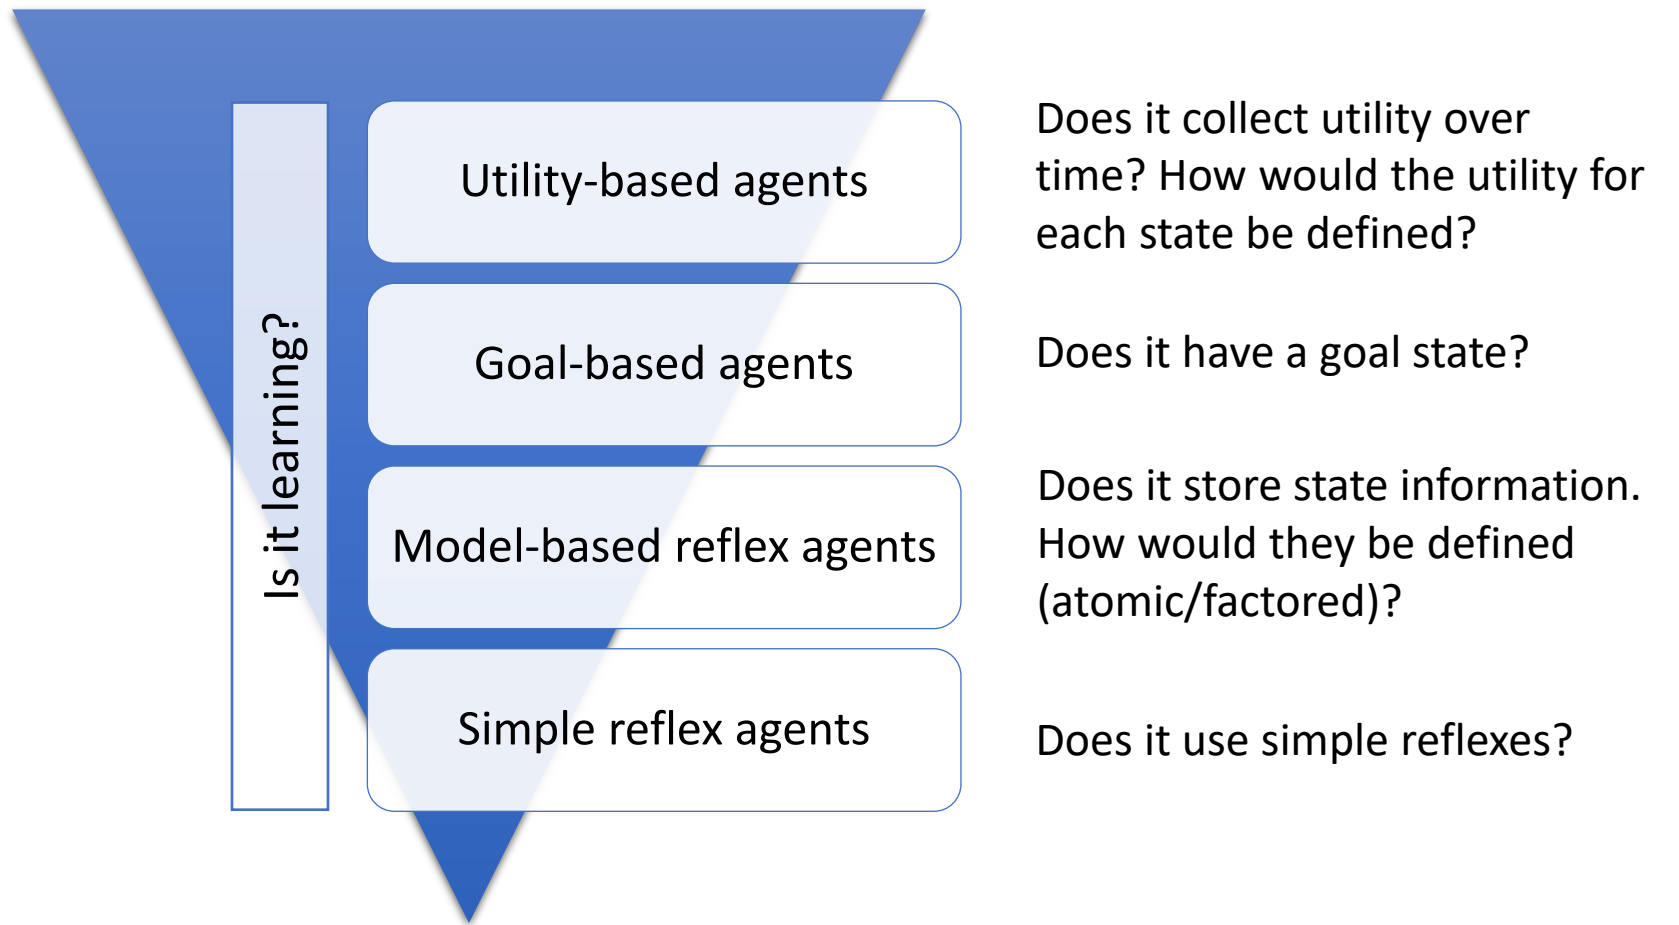
iRobot's Roomba brand has become as synonymous with robot vacuum as Q-tips is with cotton swabs. The Wi-Fi-enabled Roomba 960 is ample evidence why. It turns a tiresome chore into something you can almost look forward to. With three cleaning modes and dirt-detecting sensors, it kept all the floor surfaces in our testing immaculate, and its camera-driven navigation and mapping were superb. Its easy-to-use app provides alerts and detailed cleaning reports. The ability to control it with Amazon Alexa and Google Home voice commands are just the cherry on top.

Source: <https://www.techhive.com/article/3269782/best-robot-vacuum-cleaners.html>

# PEAS Description of a Modern Robot Vacuum

Performance measure	Environment	Actuators	Sensors

# What Type of Intelligent Agent is a Modern Robot Vacuum?



Check what applies



# What Type of Intelligent Agent is this?



Default (GPT-3.5)



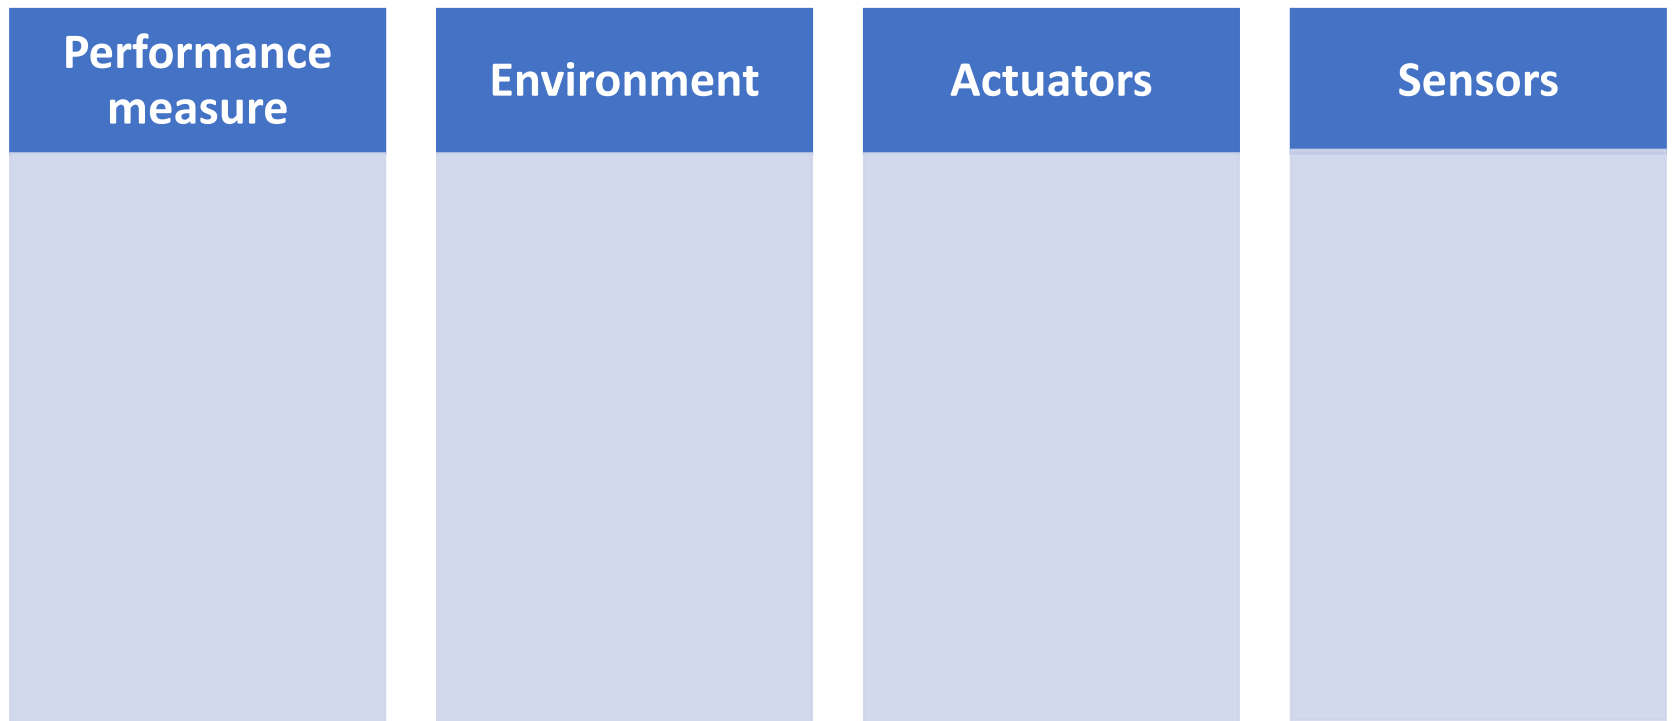
the sun is shining. It is



a beautiful day. The warm rays of sunlight illuminate the surroundings, creating a cheerful and inviting atmosphere. People are out and about, enjoying outdoor activities, and the clear blue sky adds to the overall sense of positivity. Whether you're taking a leisurely stroll, having a picnic, or simply basking in the sun's glow, it's a perfect time to appreciate the natural beauty and soak in the uplifting energy that the sunshine brings.

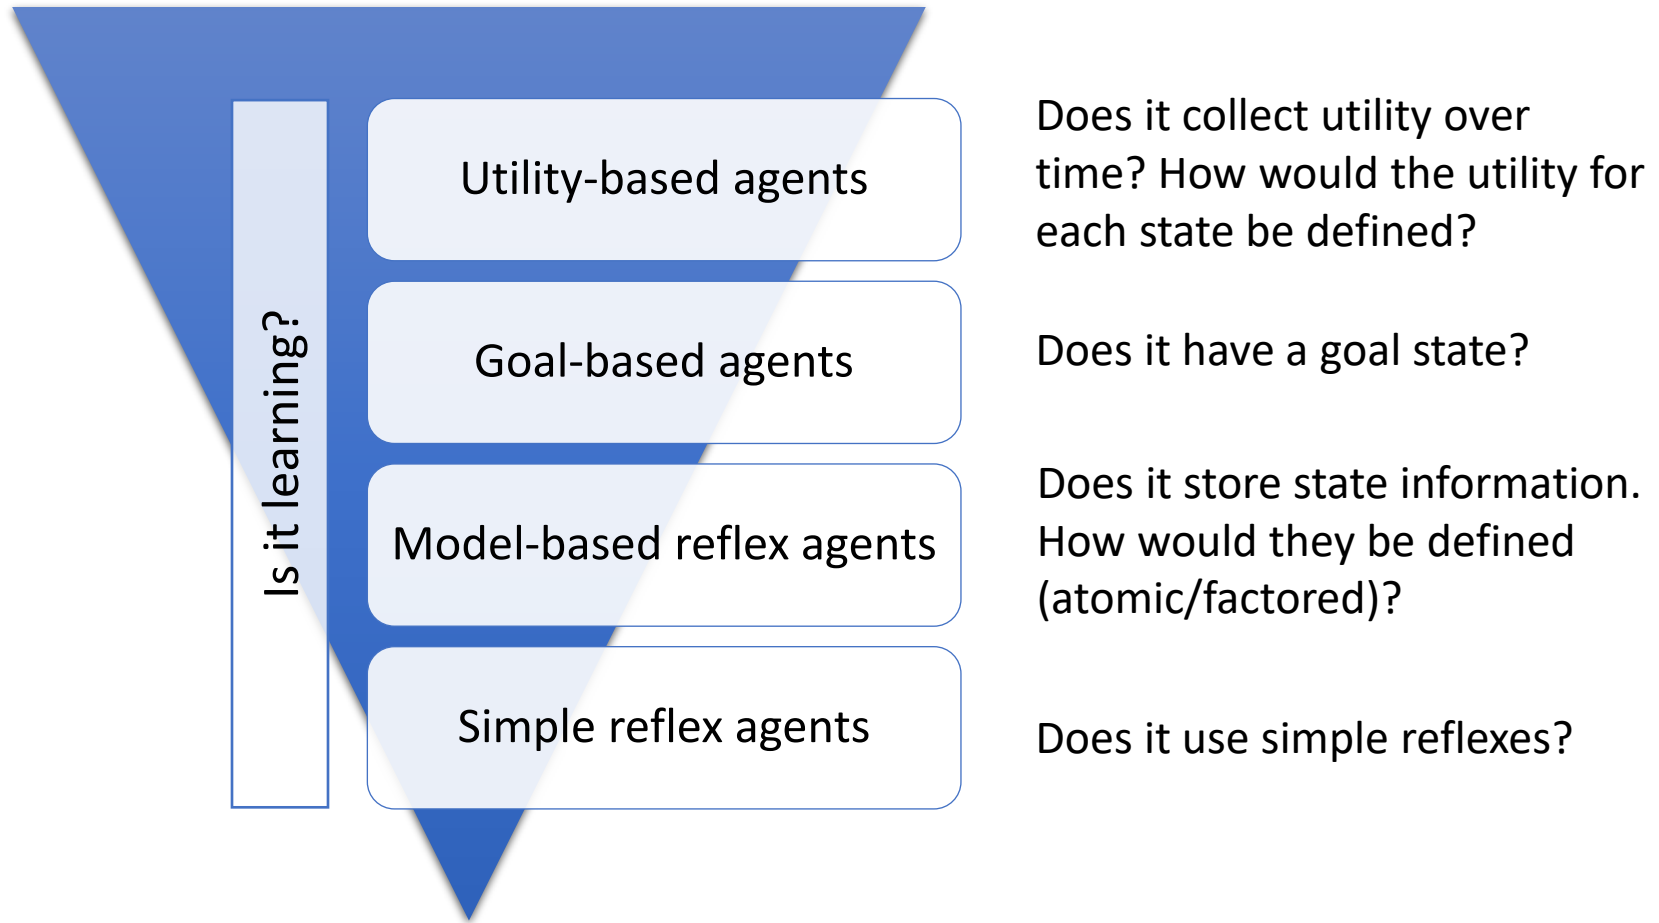


# PEAS Description of ChatGPT



How does ChatGPT work?

# What Type of Intelligent Agent is ChatGPT?



Check what applies

Answer the following questions:

- Does ChatGPT pass the Touring test?
- Is it ChatGPT a rational agent? Why?

# Intelligent Systems as Sets of Agents: Self-driving Car



It should learn!

Utility-based agents

Goal-based agents

Model-based reflex agents

Simple reflex agents

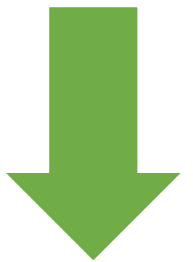
Make sure the passenger has a pleasant drive  
(not too much sudden breaking = utility)

Plan the route to the destination.

Remember where every other car is and  
calculate where they will be in the next few  
seconds.

React to unforeseen issues like a child  
running in front of the car quickly.

High-level  
planning



Low-level  
planning

# Conclusion

Intelligent agents inspire the research areas of modern AI

**Search** for a goal  
(e.g., navigation).

**Optimize** functions  
(e.g., utility).

Stay within given  
**constraints**

(constraint satisfaction problem;  
e.g., reach the goal without  
running out of power)

Deal with **uncertainty**  
(e.g., current traffic on the  
road).

**Learn** a good agent  
program from data  
and improve over time  
(machine learning).

**Sensing**  
(e.g, natural language  
processing, vision)